

PROGRAMMING IN PYTHON II

Data Analysis and Preprocessing



Michael Widrich
Institute for Machine Learning

Copyright statement:

This material, no matter whether in printed or electronic form, may be used for personal and non-commercial educational use only. Any reproduction of this material, no matter whether as a whole or in parts, no matter whether in printed or in electronic form, requires explicit prior acceptance of the authors.

Outline

1. Terminology
2. Motivation
3. Cleaning up our dataset
4. First analysis
5. Data preprocessing
6. Normalizing
7. Optimization
8. Data analysis

TERMINOLOGY



Terminology

Model: parameterized function/method with specific parameter values (e.g. a trained neural network)

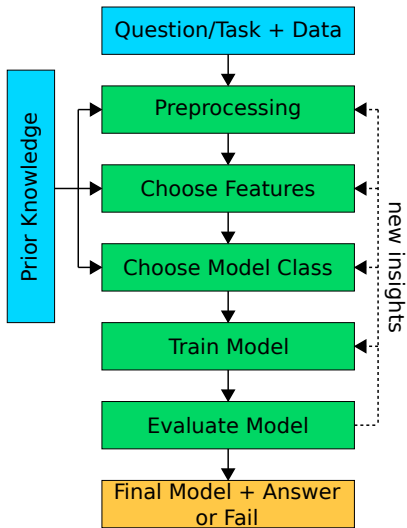
Model class: the class of models in which we search for the model (e.g. neural networks, SVMs, ...)

Parameters: representations of concrete models inside the given model class (e.g. network weights)

Hyperparameters: parameters controlling model complexity or the training procedure (e.g. network learning rate)

Model selection/training: process of finding a model from the model class

Basic data analysis workflow



MOTIVATION



Motivation

- We want to train a ML model such that we get a "good" or even the "best" model
- How do we get the "best" model?
 1. How does our model perform on our data? – Loss function
 2. How will it perform on (unseen) future data?
 - Generalization

Generalization – Theory

- Generalization is something humans hope for every day
...but sometimes fail at

Generalization – Theory

- Generalization is something humans hope for every day
...but sometimes fail at
- Generalization in a nut-shell:
 - We train our model on a subset of data points (e.g. to predict labels)
 - We use **Empirical Risk Minimization (ERM)**
 - This subset of data points is called **training set**
 - We want this trained model to also work on (e.g. correctly predict) unknown/future data
 - Problem: We might fit our parameters to noise specific to our training dataset (= **over-fitting**)

Generalization – Theory

- Generalization is something humans hope for every day
...but sometimes fail at
- Generalization in a nut-shell:
 - We train our model on a subset of data points (e.g. to predict labels)
 - We use **Empirical Risk Minimization (ERM)**
 - This subset of data points is called **training set**
 - We want this trained model to also work on (e.g. correctly predict) unknown/future data
 - Problem: We might fit our parameters to noise specific to our training dataset (= **over-fitting**)
 - We can use separate subset of samples to estimate the (true) risk on unknown data (= how well our model **generalizes**)
 - This separate subset of data points is called **test set**

Generalization – Assumptions

- Of course there is a price to pay: The theory comes with assumptions. . .
 1. Strong law of large numbers: Our subset of data points has to be large enough
 2. Our data points have to be independently and identically distributed (i.i.d.)

What does i.i.d. mean in practice? (1)

- Independently and identically distributed (i.i.d.):
Each sample has the same probability distribution as the others and all are mutually independent.

What does i.i.d. mean in practice? (2)

- Example: Our ML project

What does i.i.d. mean in practice? (2)

■ Example: Our ML project

□ What we want

- We want our model to perform data imputation on all kinds of pictures within certain restrictions (size, color)
- The distribution our pictures are sampled from should be that of all possible pictures within these restrictions
- The pictures should be sampled randomly from this distribution of all possible pictures within these restrictions

What does i.i.d. mean in practice? (2)

■ Example: Our ML project

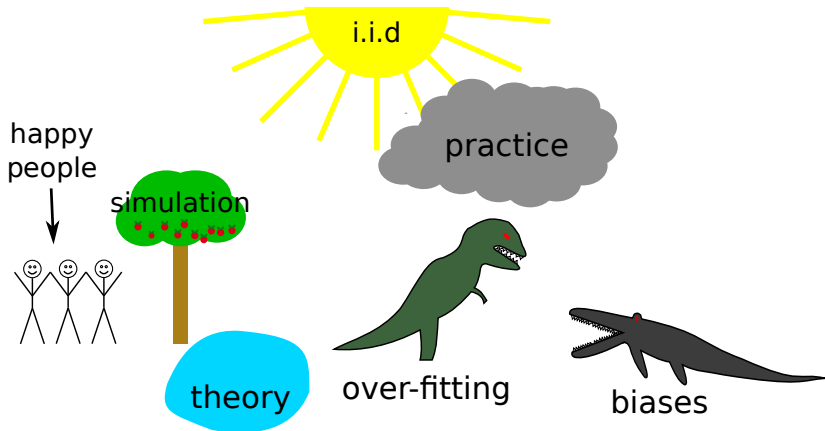
□ What we want

- We want our model to perform data imputation on all kinds of pictures within certain restrictions (size, color)
- The distribution our pictures are sampled from should be that of all possible pictures within these restrictions
- The pictures should be sampled randomly from this distribution of all possible pictures within these restrictions

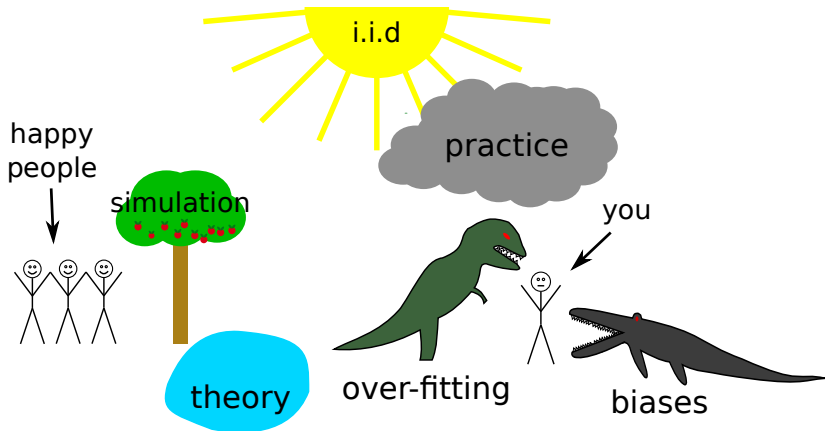
□ What we have

- We collected 100 pictures per student
- The 100 pictures per student are probably not mutually independent
- Pictures not sampled from distribution of all possible pictures but other distribution (European setting, ML-students, ...)
- Pictures are not randomly drawn from the true distribution of all possible pictures!

You and i.i.d. in practice



You and i.i.d. in practice



Working with what we have – Theory

- We need to consider the violations of i.i.d. properties in our data
- Training set and test set splitting must reflect this consideration
 - Test set must be drawn independently from training set (or as independently as possible) to get a good estimate of true risk
 - Preprocessing must not violate test and training set split
 - Data analysis done on complete set of data points cannot be used for training

Working with what we have – Practice

- Example: Our ML project

Working with what we have – Practice

■ Example: Our ML project

- Random assignment of samples to training and test set will not be sufficient! (Not independently sampled!)
- Better: Assign samples of one set of students to the training set and those of other students to the test set
- In general: Assign samples of one set of clusters to the training set and those of other clusters to the test set if you want estimate for generalization between clusters!

Working with what we have – Practice

■ Example: Our ML project

- Random assignment of samples to training and test set will not be sufficient! (Not independently sampled!)
- Better: Assign samples of one set of students to the training set and those of other students to the test set
- In general: Assign samples of one set of clusters to the training set and those of other clusters to the test set if you want estimate for generalization between clusters!

■ Even then we will not get rid of the problem that we did not sample correctly from distribution of all possible pictures

- We do not know how well our model performs on distribution of all possible pictures

Working with what we have – Practice

- Example: Our ML project

- ☐ Random assignment of samples to training and test set will not be sufficient! (Not independently sampled!)
- ☐ Better: Assign samples of one set of students to the training set and those of other students to the test set
- ☐ In general: Assign samples of one set of clusters to the training set and those of other clusters to the test set if you want estimate for generalization between clusters!

- Even then we will not get rid of the problem that we did not sample correctly from distribution of all possible pictures

- We do not know how well our model performs on distribution of all possible pictures

- Keeping that in mind, let's try our luck and get started!

CLEANING UP OUR DATASET



Cleaning up

- Never trust in that the data is valid or correctly formatted

Cleaning up

- Never trust in that the data is valid or correctly formatted
- Typical problems
 1. Empty or corrupted files
 2. Wrong filetypes
 3. Duplicated datapoints
 4. Inconsistent filenames/samplenames
 5. Inconsistent label names

Cleaning up

- Never trust in that the data is valid or correctly formatted

- Typical problems

1. Empty or corrupted files
2. Wrong filetypes
3. Duplicated datapoints
4. Inconsistent filenames/samplenames
5. Inconsistent label names

→ Exercise 2 will be on cleaning up the data

FIRST ANALYSIS



First analysis

- Check mean/standard deviation of data points
- Check number of valid samples
- Check number of classes and valid labels
- Our ML project: ~27k valid samples from ~275 students

DATA PREPROCESSING



Data preprocessing

- What violates the training and test split?
 - Do not compute global values for the whole dataset for normalization!
 - Do not perform feature-selection on the whole dataset!
- What preprocessing should be done once and saved and what should be done on-the-fly?

NORMALIZING



Normalizing

- Many ML methods profit from normalized data
 - Make data more homogeneous
 - Reduce chances to over-fit
 - Some methods require a specific normalization
- Typical for NN: Mean=0, Variance=1
- Clustering and down-projection methods also benefit from normalized data
- Normalization per sample: Mean and variance computed per sample
 - does not violate dataset splits!

Normalizing

- Many ML methods profit from normalized data
 - Make data more homogeneous
 - Reduce chances to over-fit
 - Some methods require a specific normalization
- Typical for NN: Mean=0, Variance=1
- Clustering and down-projection methods also benefit from normalized data
- Normalization per sample: Mean and variance computed per sample
 - does not violate dataset splits!
- Exercise 3 will be on normalizing the data

OPTIMIZATION



Optimization

- Prepare data such that we do not need to convert it before feeding it to our models
- Compress data set to save disk space
- Load data set in RAM if possible to decrease loading time
- Use folders to structure data files or load them into one container, e.g. hdf5
 - Max. number of files per directory, max. size per file, max. length of file paths depending on filesystem/OS

DATA ANALYSIS



Clustering and Down-Projection (1)

- After normalization, look into clustering and down-projection methods
 - Give us valuable insights in the data
 - If you use such clusters to create test and training split, verify them manually! (Do not trust clustering methods.)
- Popular methods: PCA, ICA, t-SNE, UMAP, ...

Clustering and Down-Projection (1)

- After normalization, look into clustering and down-projection methods
 - Give us valuable insights in the data
 - If you use such clusters to create test and training split, verify them manually! (Do not trust clustering methods.)
- Popular methods: PCA, ICA, t-SNE, UMAP, ...
- Our raw data might be incompatible with these methods
 - Too many feature values
 - Datapoints in odd feature-space
 - We need to be creative

Clustering and Down-Projection (2)

Example: Our ML project

- We would like to use t-SNE or UMAP for clustering
- What we want:
 - ☐ Small suitable feature-space
 - ☐ Constant number of features
- What we have:
 - ☐ Huge feature space (number of pixels)
 - ☐ Odd feature-space (pixel-space)
 - ☐ Images of different size (different number of features)

Clustering and Down-Projection (3)

■ Possible solution:

- Down-project images into better feature-space before clustering using pretrained CNN features
- Less features, constant number of features, better feature-space (hopefully)
- Alternative: PCA or ICA

Outlook

- Next time: Implementing and optimizing data loading using PyTorch