

# PROGRAMMING IN PYTHON II

## Neural Network Implementation: Inference



Michael Widrich  
Institute for Machine Learning

## Copyright statement:

This material, no matter whether in printed or electronic form, may be used for personal and non-commercial educational use only. Any reproduction of this material, no matter whether as a whole or in parts, no matter whether in printed or in electronic form, requires explicit prior acceptance of the authors.

# Outline

- 1. Motivation**
- 2. PyTorch for building neural networks**
- 3. Quick guide: Designing your neural network**
- 4. Python II Project**
- 5. Further reading**

# MOTIVATION



# Motivation

- In order to design our neural networks efficiently, we would like to have
  - modular structure of layers and building blocks
  - consistent interfaces between modules
  - free combination of modules (not only sequential but also tree-like structures)
  - automated optimization of computations
  - automated deployment to CPU, GPU, or other dedicated hardware
  - automated computation of gradients for training

# Motivation

- In order to design our neural networks efficiently, we would like to have
    - modular structure of layers and building blocks
    - consistent interfaces between modules
    - free combination of modules (not only sequential but also tree-like structures)
    - automated optimization of computations
    - automated deployment to CPU, GPU, or other dedicated hardware
    - automated computation of gradients for training
- PyTorch provides these features!

# PYTORCH FOR BUILDING NEURAL NETWORKS



# torch.nn.Module

- torch.nn.Module is the base class for all neural network modules
- All custom networks/layers/modules should be derived from this class
- Can contain and utilize other modules
- Automated registration of trainable parameters and parameters in submodules



# torch.nn.Module: Usage

- torch.nn.Module usage is simple:
  1. Create class that inherits from torch.nn.Module
  2. Define a `.forward()` method
  3. Overwrite the `.init()` method
  4. Create an instance of your class and apply it to input

## torch.nn.Module: Details

- The `.forward()` method will be executed when your class instance is applied to an input
- Using PyTorch (sub)module instances as module attributes, will register the submodule automatically
- Using PyTorch parameter instances as module attributes, will register the parameter automatically
  - `torch.nn.Parameter` will create trainable tensors, e.g. for NN weights
- PyTorch modules behave a lot like PyTorch tensors
  - Can be sent to devices using `.to(device=...)`
  - Can be converted to datatype using `.to(dtype=...)`

# Predifined PyTorch modules

- Many predefined PyTorch modules exist
- Sometimes include additional optimization
  - E.g. PyTorch LSTM with specialized CUDA support but less flexible design
- Should be preferred over custom modules unless special functionality is desired

# QUICK GUIDE: DESIGNING YOUR NEURAL NETWORK



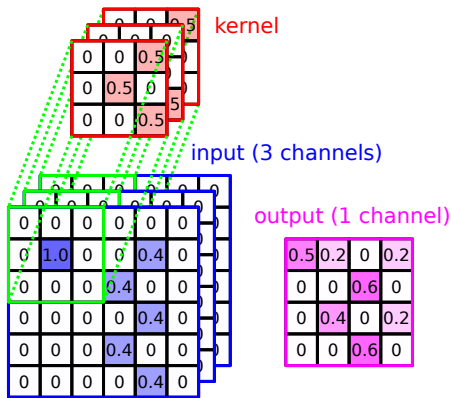
# NN types: FFNN

- Fully-connected feed-forward NN (FFNN):
  - High complexity, high number of weights
  - Not utilizing order in feature vectors
  - Multi-dimensional inputs are typically flattened (reshaped to flat vector as input)

# NN types: CNN (1)

## ■ Convolutional NN (CNN):

- Weight kernels are slid along (=convolved) an input tensor



## NN types: CNN (2)

### ■ Convolutional NN (CNN):

- 1D CNNs: Kernels convolved over one dimension (e.g. time dimension in a time-series)
- 2D CNNs: Kernels convolved over 2 dimensions (e.g. spatial dimensions in image)
- nD CNNs: Possible but uncommon at the moment (expensive/lacking optimization)
- Less complexity, fixed kernel size to incorporate information about adjacent features
- Depth of network increases field of vision
- Typically employ pooling operations to pool adjacent features (e.g. max-pooling)
- Outputs sometimes flattend and then fed into FFNN or max-pooled to 1D feature vector

# NN types: RNN

## ■ Recurrent NN (RNN):

- ☐ NN applied to a list of inputs, using the same weights for each input
- ☐ NN has access to output/hidden state from previous input (=recurrent)
- ☐ Turing complete (but need to be trained somehow)
- ☐ 1D: Sequence (=list of feature vectors) as input
- ☐ nD: Multidimensional RNN variants for matrices or tree-like structure
- ☐ Most popular, since no **vanishing gradient**: LSTM
- ☐ Alternative to RNNs: Transformer NNs

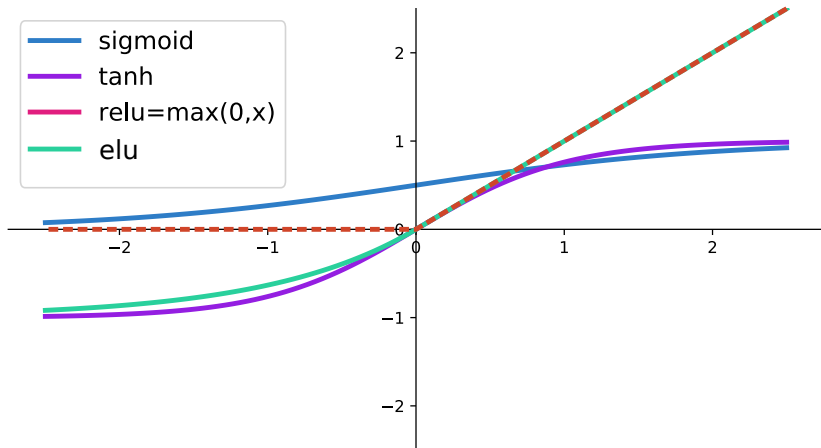


# Common activation functions

- Pay attention to recommended weight initialization and input normalization for the specific activation functions
- Scaled Exponential Linear Unit (SELU)
  - Self-normalizing properties benefit learning and design
  - Common for FFNN but also successes in CNN
- Rectified Linear Unit (ReLU)
  - Very common in CNNs
  - Units can “die” once they get stuck in 0-activation (=no gradients)
- Sigmoid activation function
  - Common in (older) NNs
  - Introduce vanishing gradients
  - Outputs in range  $[0, 1]$

# Common activation functions

Activation functions



# PYTHON II PROJECT



# Python II Project: NN Design

- We deal with image data → CNN is a natural choice
- Hyperparameters (see Unit 07): Number of kernels, size of kernels, pooling operations, activation function, skip connections, and number of layers
- Additional data can be fed to CNN by creating a new feature channel (see Assignment 2)
- CNN could be combined with FFNN but probably not necessary
- SELU or ReLU activation functions are good candidates

## FURTHER READING



## Further reading

- ÖAW AI summer school slide-set:  
[https://github.com/ml-jku/oeaw\\_ai\\_summer\\_school](https://github.com/ml-jku/oeaw_ai_summer_school)
- Courses and lecture materials in AI-study (Hands-on AI, Machine Learning: Supervised Techniques, LSTM and Recurrent Neural Nets, ...)
- *Pattern Recognition and Machine Learning* (C. Bishop)
- *Dive into Deep Learning* (A. Zhang, Z. Lipton, M. Li, A. Smola): <https://d2l.ai/>
- <https://pytorch.org/tutorials/index.html>
- [https://pytorch.org/tutorials/beginner/blitz/neural\\_networks\\_tutorial.html](https://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html)